Functional Modeling with Neural Causal Models and Personalized Treatment Effect Estimation

Mike Krähenbühl, Supervisors: Beate Sick, Oliver Dürr

July 26, 2025

Background

Supervisors:

— Beate Sick, UZH

— Oliver Dürr, HTWG Konstanz

Paper *"Interpretable Neural Causal Models with TRAM-DAGs"* (Sick and Dürr, 2025):

— Framework to model causal relationships

— Based on transformation models

— Rely on (deep) neural networks

— Compromise between interpretability and flexibility

They showed on synthetic data, that TRAM-DAGs can be fitted on observational data and tackle causal queries on all three levels of Pearl's causal hierarchy.

Research Questions

In this presentation:

1. TRAM-DAGs
   — How do they work?
2. Individualized Treatment Effect (ITE) estimation
   — Does it work on real data (International Stroke Trial)?
   — When and why does ITE estimation fail (simulation)?
   — How to estimate ITEs with TRAM-DAGs in a complicated graph (simulation)?
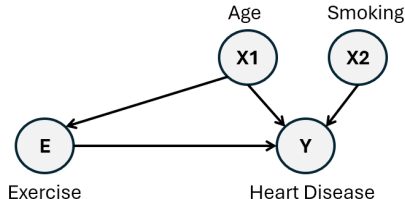
TRAM-DAGs

TRAM-DAGs: Motivation

Randomized Controlled Trial:

— Gold standard for estimating causal effect

— Solves problem of confounding
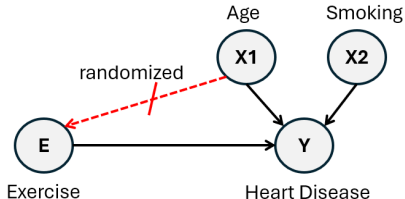
Observational Data:

— Real world, potential confounding

— We assume no unobserved confounding

TRAM-DAGs: Motivation

Pearl's causal hierarchy (Pearl, 2009)

Observational: $P(Y = 1 \mid E = 1)$
"Probability of heart disease given that the person exercises"

Interventional: $P(Y = 1 \mid \text{do}(E = 1))$
"Probability of heart disease if we made people start exercising"

Counterfactual: $P(Y_{(E=1)} = 1 \mid E = 0, Y = 1)$
"Would someone who does not exercise and has heart disease still have it if they had exercised?"

TRAM-DAGs: Background

Structural Causal Model: Describes the causal mechanism and probabilistic uncertainty (Pearl, 2009)

— X_i = observed variable
— Z_i = noise distribution
— f_i = deterministic function: $X_i = f_i(Z_i, \text{pa}_i)$

$\rightarrow$ We want a model that estimates $X_i = f_i(Z_i, \text{pa}_i)$ in a flexible and interpretable way!

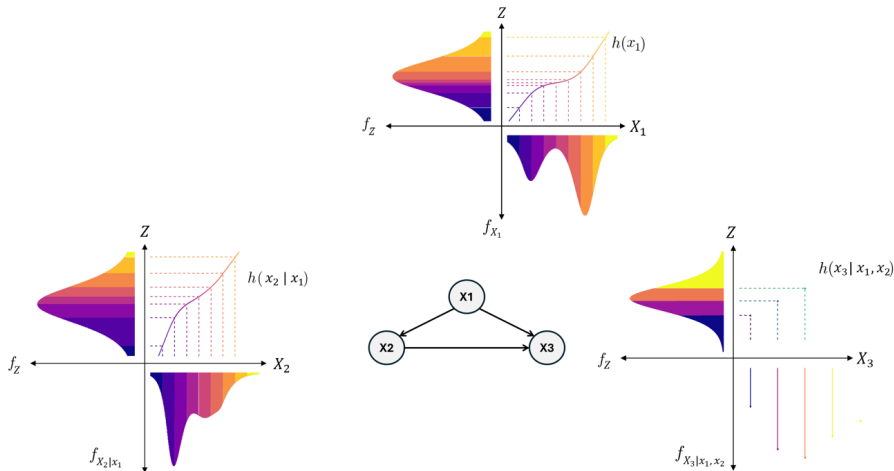$$Z \sim F_{Z_1}, Z_2 \sim F_{Z_2}, Z_3 \sim F_{Z_3}$$

$$X_1 = f_1(Z_1)$$
$$X_2 = f_2(Z_2, X_1)$$
$$X_3 = f_3(Z_3, X_1, X_2)$$

TRAM-DAGs: Background

Proposed framework: TRAM-DAGs (Sick and Dürr, 2025)

TRAM-DAGs: Background

Transformation Models: Flexible distributional regression method
(Hothorn et al., 2014)

Continuous $Y \in \mathbb{R}$:

$$F_{Y|\mathbf{X}=\mathbf{x}}(y) = F_Z(h(y) + \mathbf{x}^\top \boldsymbol{\beta})$$

Discrete $Y \in \{y_1, y_2, \ldots, y_K\}$:

$$P(Y \leq y_k \mid \mathbf{X} = \mathbf{x}) = F_Z(\vartheta_k + \mathbf{x}^\top \boldsymbol{\beta}), \quad k = 1, 2, \ldots, K-1$$

— F_Z: CDF of the latent distribution (e.g. standard logistic)
— h: Transformation function, monotonically increasing
— $\mathbf{x}$: Predictors

TRAM-DAGs: Background

Extended to Deep TRAMs (Sick et al., 2021)

— Customizable transformation model using neural networks (NNs)
— Minimizing negative log-likelihood (NLL) via NN optimization

Effects of predictors: LS (Linear Shift), CS (Complex Shift), CI (Complex Intercept)

$$h(y|\boldsymbol{x_L}, \boldsymbol{x_C}) = \text{SI} + \text{LS}_{\boldsymbol{x_L}} + \text{CS}_{\boldsymbol{x_C}}$$

TRAM-DAGs: Experiment 1 (simulation)

Setup:

— We have:
  — Observational data (simulated)
  — Predefined DAG
— We want:
  — Estimate $Z_i = h_i(X_i \mid \mathrm{pa}(X_i))$ of each variable i
  — Sample from conditional distributions for causal queries with structural equations $X_i = h_i^{-1}(Z_i \mid \mathrm{pa}(X_i))$


$X_1 \sim F_z(h(x_1))$

$X_2 \sim F_z(h(x_2) + \mathrm{LS}_{x1})$

$X_3 \sim F_z(h(x_3) + \mathrm{LS}_{x1} + \mathrm{CS}_{x2})$

TRAM-DAGs: Experiment 1 (simulation)

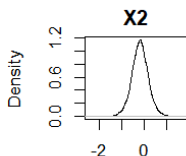**Data-generating process (DGP):**

X_1**:** Continuous, bimodal. *Source node* (independent).

X_2**:** Continuous. Depends on X_1 (linear):

$$\beta_{12} = 2, \quad h_l(X_2) = 5X_2$$

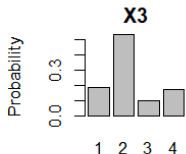$$\boxed{h(X_2 \mid X_1) = h_l(X_2) + \beta_{12}X_1}$$

X_3**:** Ordinal. Depends on X_1 (linear) and X_2 (complex):

$$\beta_{13} = 0.2, \quad f(X_2) = 0.5 \cdot \exp(X_2), \quad \vartheta_k \in \{-2,\ 0.42,\ 1.02\}$$

$$\boxed{h(X_{3,k} \mid X_1, X_2) = \vartheta_k + \beta_{13}X_1 + f(X_2)}$$

X1

X2

X3

TRAM-DAGs: Experiment 1 (simulation)

Construct Model: Modular Neural Network

Inputs: Observations + assumed structure

Outputs:

— Simple Intercepts (SI): ϑ

— Linear Shifts (LS): $\beta_{12}X_1$, $\beta_{13}X_2$

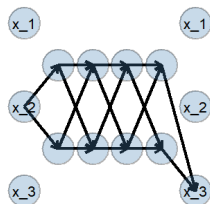— Complex Shift (CS): $f(X_2)$

Assemble transformation functions:

$$\boxed{h(X_i \mid \text{pa}(X_i)) = \text{SI} + \text{LS} + \text{CS}}$$

$$h(X_1) = h_I(X_1)$$

$$h(X_2 \mid X_1) = h_I(X_2) + \beta_{12}X_1$$

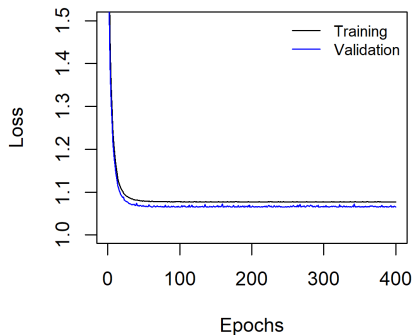$$h(X_{3,k} \mid X_1, X_2) = \vartheta_k + \beta_{13}X_1 + f(X_2)$$

CS_{X_2} on X_3

TRAM-DAGs: Experiment 1 (simulation)

Model fitting: 20,000 training samples, 400 epochs

Sampling from the Fitted TRAM-DAG

Nodes $X_i, i \in \{1, 2, 3\}$**:**

— Sample latent value:

$$z_i \sim F_{Z_i} \quad (\text{e.g., } \texttt{rlogis()} \text{ in R})$$

— Determine x_i such that:
  — **If X_i is continuous:** Solve for x_i using numerical root-finding:
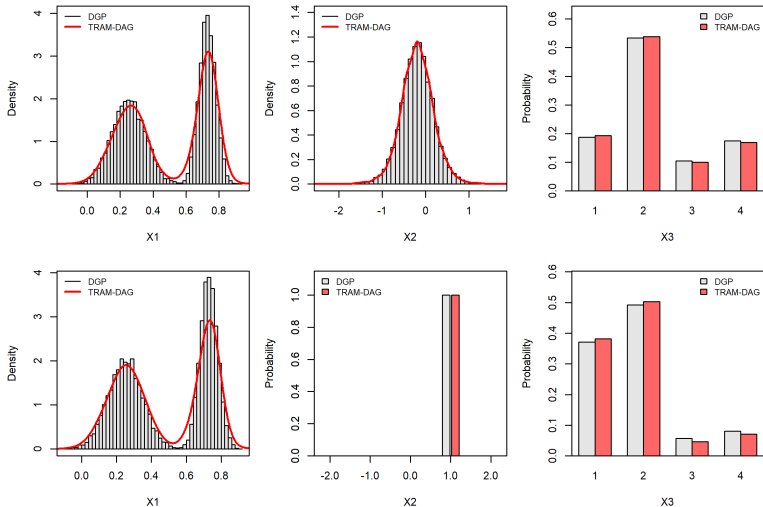
$$h(x_i \mid \text{pa}(x_i)) - z_i = 0$$

  — **If X_i is ordinal:** find the smallest category x_i such that

$$x_i = \max\left(\{0\} \cup \{x : z_i > h(x \mid \text{pa}(x_i))\}\right) + 1$$

TRAM-DAGs: Experiment 1 (simulation)

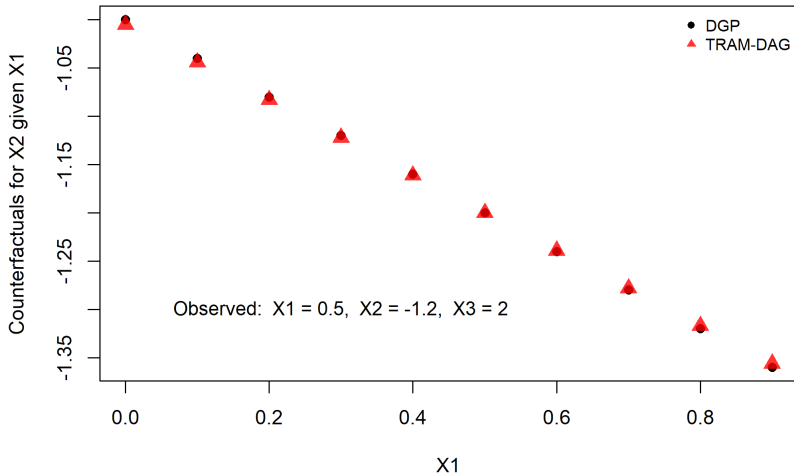## Sampled observational and interventional distributions:

Experiment 1: TRAM-DAGs (simulation)

How to determine a counterfactual value for X_2, given some observation?

Experiment 1: TRAM-DAGs (simulation)

Counterfactuals: Counterfactual value of X_2 under varying X_1

Experiment 1: TRAM-DAGs (simulation)

Discussion: With TRAM-DAGs we can

— estimate the functional form of the edges in the DAG
— customize flexibility and interpretability (SI/CI, LS, CS)
— sample from the fitted model (observational/interventional)
— estimate counterfactuals

Individualized Treatment Effects (ITEs)

Individualized Treatment Effect (ITE): Motivation

Why ITE?

— RCTs estimate the Average Treatment Effect (ATE)

— Individuals may respond differently based on covariates

— Important for personalized medicine, targeted marketing, etc.

— Heterogeneity mostly driven by treatment-covariate interactions

Definition: Difference in potential outcomes (Rubin, 2005)

$$Y_i(1) - Y_i(0)$$

where $Y_i(1)$: outcome if treated, $Y_i(0)$: if not treated

Fundamental problem: We never observe both $Y_i(1)$ and $Y_i(0)$ for the same individual (Holland, 1986)

From Unobservable to Estimable ITE

Goal: Estimate the *individualized treatment effect (ITE)* from observed data (Hoogland et al., 2021).

$$\text{ITE}(\mathbf{x}_i) = \mathbb{E}[Y_i(1) - Y_i(0) \mid \mathbf{X} = \mathbf{x}_i]$$

$$= \mathbb{E}[Y_i(1) \mid T = 1, \mathbf{X} = \mathbf{x}_i] - \mathbb{E}[Y_i(0) \mid T = 0, \mathbf{X} = \mathbf{x}_i]$$

(by ignorability/exchangeability: no unmeasured confounding)

$$= \mathbb{E}[Y_i \mid T = 1, \mathbf{X} = \mathbf{x}_i] - \mathbb{E}[Y_i \mid T = 0, \mathbf{X} = \mathbf{x}_i]$$

(by consistency: observed = potential outcome, e.g. correct label)

Further assumptions:

— **Positivity:** every individual could receive either treatment (e.g. no deterministic assignment)

— **No interference:** one person's treatment doesn't affect another's outcome

Individualized Treatment Effect (ITE): Models

How did we estimate the potential outcomes $\mathbb{E}[Y_i \mid T = t, \mathbf{X} = \mathbf{x}_i]$?

— **T-learner:**

  1. Fit two separate models on treated and control groups
  2. Predict $\mathbb{E}[Y_i \mid \mathbf{X} = \mathbf{x}_i]$ from each model

  — Logistic regression / Random forest (with hyperparameter tuning)

— **S-learner:**

  1. Fit one model on all data with treatment as a feature
  2. Predict $\mathbb{E}[Y_i \mid \mathrm{do}(T = t), \mathbf{X} = \mathbf{x}_i]$ by setting $T = 0$ and $T = 1$

  — TRAM-DAGs (SCM, flexible, interactions, generative)

Experiment 2: ITE on International Stroke Trial (IST)

Background/Motivation: Chen et al. (2025) showed that results of models used for ITE estimation did not generalize to the test set.

International Stroke Trial (IST):

— Large RCT on stroke patients (19,435 patients, 21 baseline covariates)
— Evaluated the effects of aspirin on stroke patients
— Binary treatment and outcome

Research question: Do we reach similar conclusion as Chen et al. (2025) when estimating ITEs with T-learners (logistic regression, tuned random forest) and S-learner (TRAM-DAGs) on IST dataset.

Experiment 2: ITE on International Stroke Trial (IST)

Results: with T-learner tuned random forest (`comets` package ([Kook, 2024](#))):
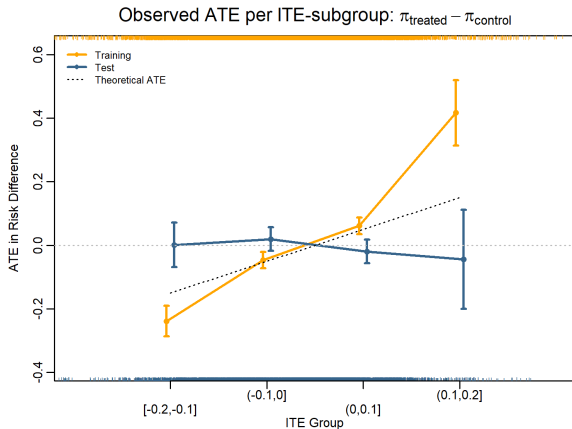
Experiment 2: ITE on International Stroke Trial (IST)

Discussion:

— We obtained similar results as Chen et al. (2025)

— Some models suggest a range of ITEs, but these ITEs do not generalize
  to the test set (no effect)

— We do not know why, since ground truth is unknown

Experiment 3: ITE model robustness in RCTs (simulation)

Motivation: ITE estimation failed on the real-world RCT of the International Stroke Trial (IST). We want to know why.

Research question: What factors contribute to the failure of ITE estimation in causal models?

Setup:

— Simulate different RCT scenarios to understand when ITE estimation fails

— Apply simple model (logistic regression; matching DGP) and non-parametric model (tuned random forest)

Simulation Case 1: Fully Observed

Setup:

— $n = 20{,}000$

— $T \sim \text{Bernoulli}(0.5)$

— $\mathbf{X} = (X_1, \ldots, X_5)^\top \sim \mathcal{N}(\mathbf{0}, \Sigma)$

— $\mathbf{X_{TX}} = (X_1, X_2)^\top$ interaction

Outcome model:

$$\mathbb{P}(Y = 1 \mid \mathbf{X}, T) = \text{logit}^{-1}\left(\beta_0 + \beta_T T + \boldsymbol{\beta}_X^\top \mathbf{X} + T \cdot \boldsymbol{\beta}_{TX}^\top \mathbf{X_{TX}} \right)$$

Simulation Case 1: Fully Observed

Results with T-learner logistic regression (glm):

Simulation Case 1: Fully Observed

Results with T-learner tuned random forest (`comets` package):
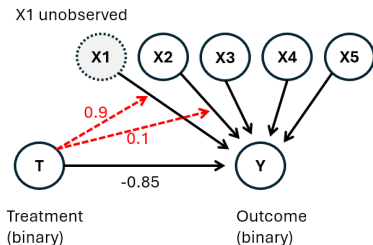
Simulation Case 1: Fully Observed

Results with (untuned) T-learner random forest (`randomForest` package):

Simulation Case 2: Unobserved Interaction

Setup:

- $n = 20{,}000$
- $T \sim \text{Bernoulli}(0.5)$
- $\mathbf{X} = (X_1, \ldots, X_5)^\top \sim \mathcal{N}(\mathbf{0}, \Sigma)$
- $\mathbf{X_{TX}} = (X_1, X_2)^\top$ interaction

X1 unobserved

0.9
0.1
-0.85

Treatment
(binary)

Outcome
(binary)

Outcome model:

$$\mathbb{P}(Y = 1 \mid \mathbf{X}, T) = \text{logit}^{-1}\left(\beta_0 + \beta_T T + \boldsymbol{\beta}_X^\top \mathbf{X} + T \cdot \boldsymbol{\beta}_{TX}^\top \mathbf{X_{TX}} \right)$$

Note: Same DGP, but X_1 is not observed!

Simulation Case 2: Unobserved Interaction

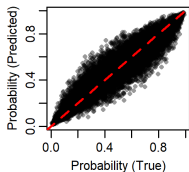Results with T-learner logistic regression (glm):

# Simulation Case 2: Unobserved Interaction

Results with T-learner tuned random forest (`comets` package):

Simulation Case 3: Fully Observed, Small Effects

Setup:

— $n = 20{,}000$

— $T \sim \text{Bernoulli}(0.5)$

— $\mathbf{X} = (X_1, \ldots, X_5)^\top \sim \mathcal{N}(\mathbf{0}, \Sigma)$

— $\mathbf{X_{TX}} = (X_1, X_2)^\top$ interaction

Treatment (binary)

Outcome (binary)

Outcome model:

$$\mathbb{P}(Y = 1 \mid \mathbf{X}, T) = \text{logit}^{-1}\left(\beta_0 + \beta_T T + \boldsymbol{\beta}_X^\top \mathbf{X} + T \cdot \boldsymbol{\beta}_{TX}^\top \mathbf{X_{TX}}\right)$$
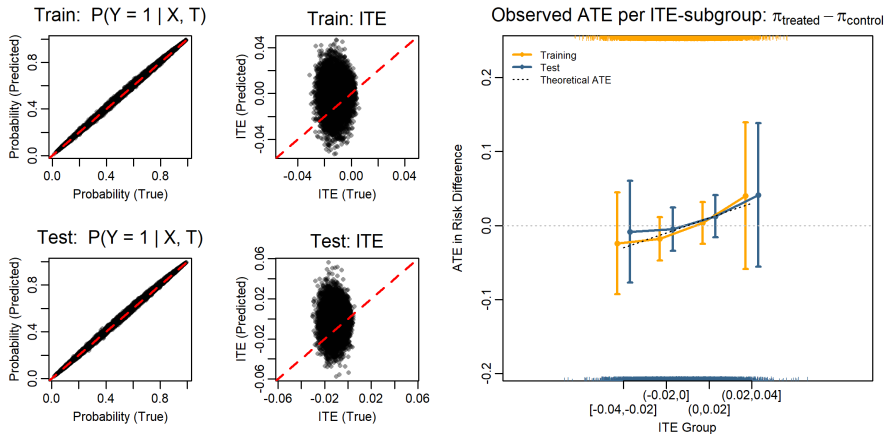
Note: Same DGP, but weak treatment effects!

# Simulation Case 3: Fully Observed, Small Effects

Results with T-learner logistic regression (glm):

Simulation Case 3: Fully Observed, Small Effects

Results with T-learner Random Forest (`comets` package):

Experiment 3: ITE model robustness in RCTs (simulation)

Key Insights:

— **Calibration** and tuning of models are crucial for reliable ITE estimation

— Ignorability alone may not guarantee unbiased ITEs if important **effect modifiers are unobserved**

— **Low true heterogeneity** may be mistaken for model failure

These factors may explain the limited ITE performance in the IST dataset.

TRAM-DAGs: Example for ITE estimation

DGP:

- $X5 = h_5^{-1}(\epsilon - 0.8\,X4) \quad \rightarrow$ (depends on treatment)
- $X6 = h_6^{-1}(\epsilon + 0.5\,X5) \quad \rightarrow$ (depends on treatment through X5)
- $Y = h_7^{-1}(\epsilon - \beta_1 X1 - \beta_2 X2 - \beta_3 X3 - \beta_4 X4 - \beta_5 X5 - \beta_6 X6 - Tr \cdot (\beta_{2Tr} X2 + \beta_{3Tr} X3))$

TRAM-DAGs: Example for ITE estimation

$$\text{ITE} = \text{median}(Y \mid \text{do}(T = 1), X) - \text{median}(Y \mid \text{do}(T = 0), X)$$

TRAM-DAGs: Example for ITE estimation

$$\text{ITE} = \text{median}(Y \mid \text{do}(T = 1), X) - \text{median}(Y \mid \text{do}(T = 0), X)$$

TRAM-DAGs: Estimate Potential Outcomes

1. Estimate each $h_i(X_i \mid \text{pa}(X_i))$ fully flexible (deep-NN / complex intercept)
2. Take the train set or a test set
3. $Z_i = h(X_i \mid pa(X_i))$ gives us the (observed) latent variable for each X_i
4. Determine counterfactuals for X5 and X6 with the (observed) latent variables Z_i
5. Determine medians of potential outcomes $Y(1)$ and $Y(0)$
6. $\text{ITE} = \text{median}(Y(1) \mid X_{tx}) - \text{median}(Y(0) \mid X_{ct})$

TRAM-DAGs: Example for ITE estimation (Results)
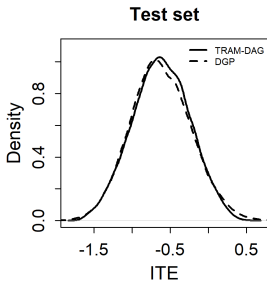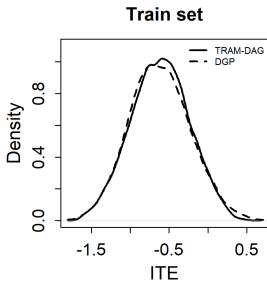
Key Findings

Findings: TRAM-DAGs

— Customizable; accurately recovers causal relationships in known DAG; allows sampling of L1-L3

— Can model interactions between variables

Findings: Individualized treatment effects (ITE)

— Calibration is important for ITE prediction

— Missing effect modifiers (or weak heterogeneity) are problematic

— TRAM-DAGs yield unbiased ITEs when DAG is correct and heterogeneity exists

Outlook

Limitations

— Simulations may not reflect real-world complexity

— TRAM-DAGs are computationally expensive (long training time)

— TRAM-DAGs require correct model specification for interpretability

— ITE estimation used median QTE for continuous outcomes

Recommendations

— Apply TRAM-DAGs to real-world datasets, including semi-structured data

— Investigate ITE estimation under unobserved effect modifiers

References I

Chen, H., Aebersold, H., Puhan, M. A., and Serra-Burriel, M. (2025). Causal machine learning methods for estimating personalised treatment effects – insights on validity from two large trials. arXiv preprint 2501.04061.

Holland, P. W. (1986). Statistics and causal inference. *Journal of the American Statistical Association*, 81(396):945–960.

Hoogland, J., IntHout, J., Belias, M., Rovers, M. M., Riley, R. D., E. Harrell Jr, F., Moons, K. G. M., Debray, T. P. A., and Reitsma, J. B. (2021). A tutorial on individualized treatment effect prediction from randomized trials with a binary endpoint. *Statistics in Medicine*, 40(26):5961–5981.

Hothorn, T., Kneib, T., and Bühlmann, P. (2014). Conditional transformation models. *Journal of the Royal Statistical Society. Series B (Statistical Methodology)*, 76(1):3–27.

Kook, L. (2024). *comets: Covariance Measure Tests for Conditional Independence*. R package version 0.1-1.

References II

Pearl, J. (2009). *Causality: Models, Reasoning and Inference*. Cambridge University Press, USA, 2nd edition.

Rubin, D. B. (2005). Causal inference using potential outcomes: Design, modeling, decisions. *Journal of the American Statistical Association*, 100(469):322–331.

Sick, B. and Dürr, O. (2025). Interpretable neural causal models with TRAM-DAGs. arXiv preprint 2503.16206, accepted at the CLeaR 2025 Conference.

Sick, B., Hothorn, T., and Dürr, O. (2021). Deep transformation models: Tackling complex regression problems with neural network based transformation models. In *2020 25th International Conference on Pattern Recognition (ICPR)*, pages 2476–2481.